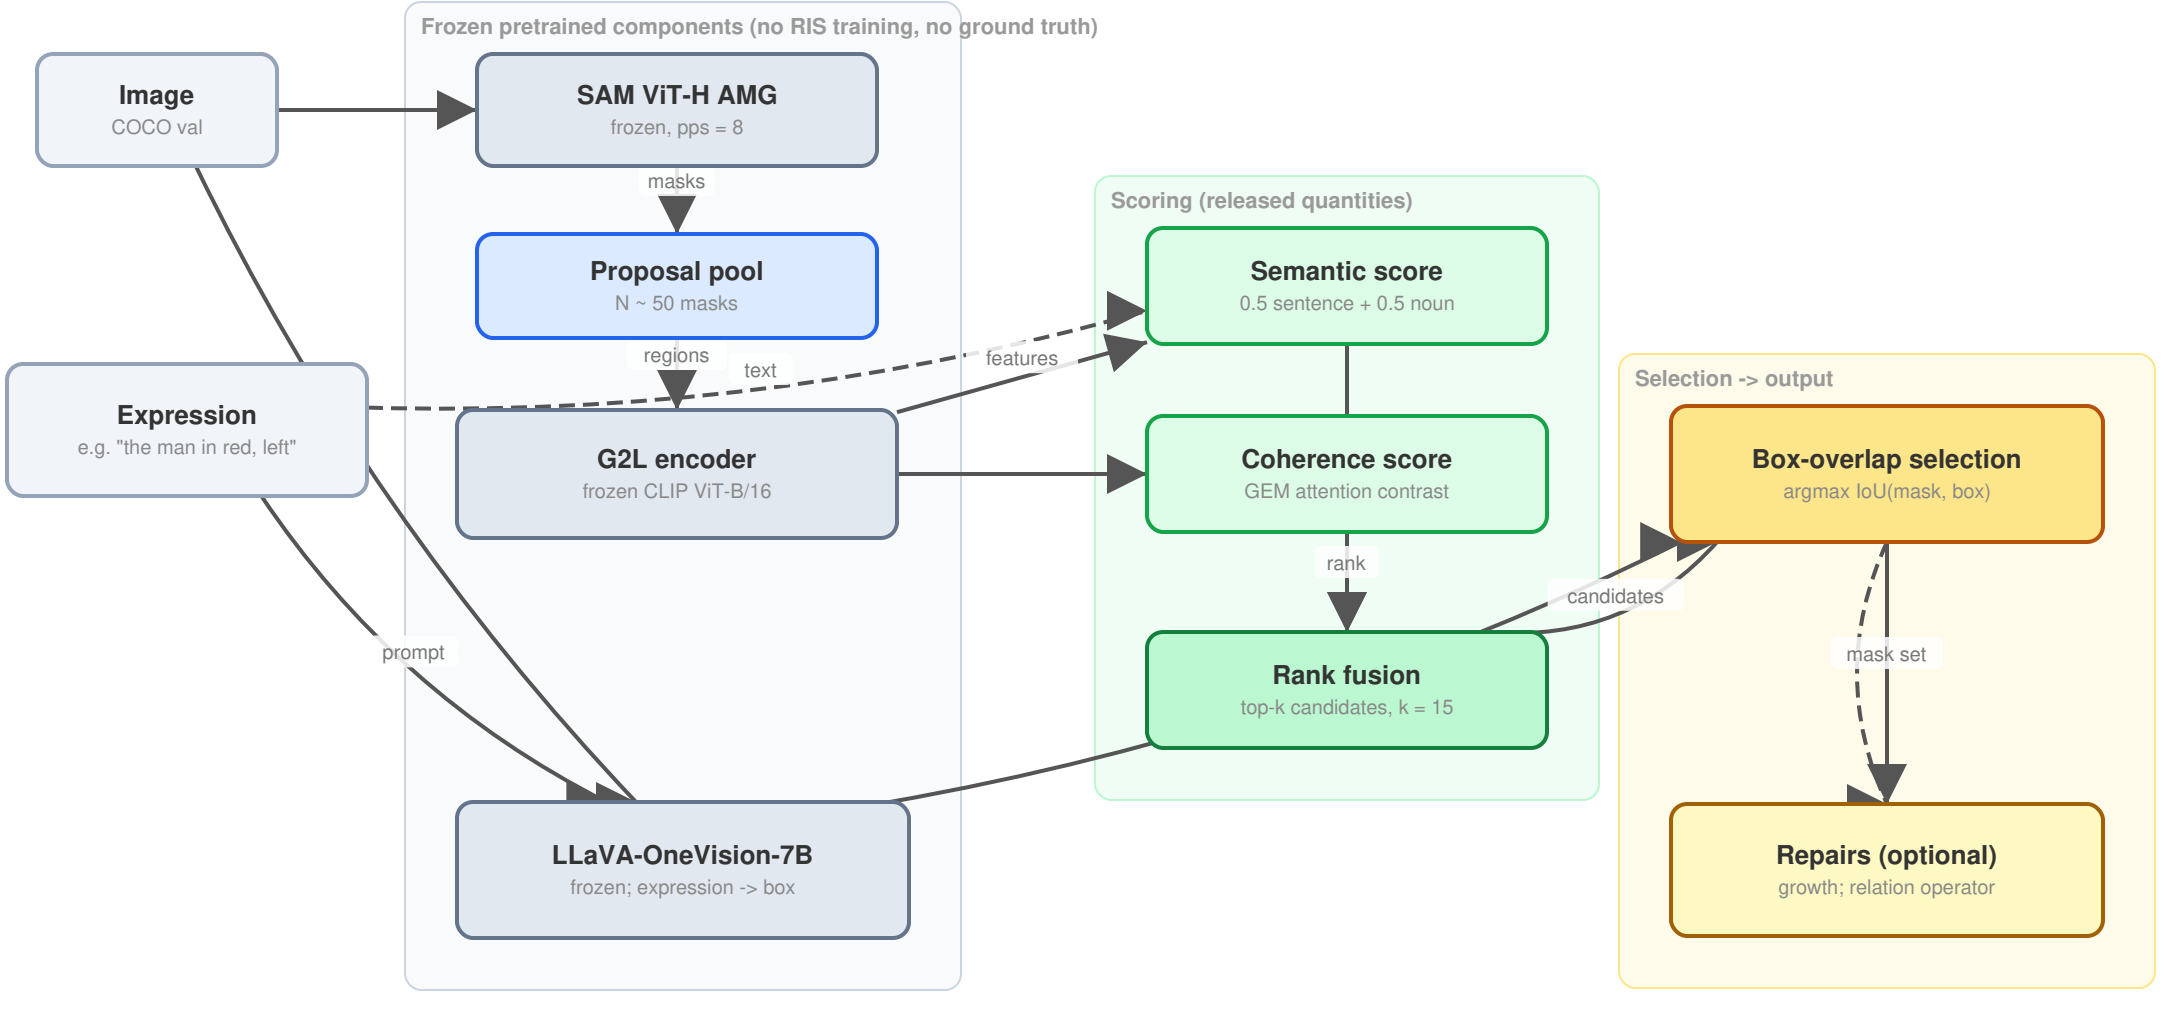
Frozen pretrained components (no RIS training, no ground truth)
Image
COCO val
SAM ViT-H AMG
frozen, pps = 8
masks
Proposal pool
N ~ 50 masks
regions
text
features
Expression
e.g. "the man in red, left"
G2L encoder
frozen CLIP ViT-B/16
prompt
LLaVA-OneVision-7B
frozen; expression -> box
Scoring (released quantities)
Semantic score
0.5 sentence + 0.5 noun
Coherence score
GEM attention contrast
rank
Rank fusion
top-k candidates, k = 15
candidates
Selection -> output
Box-overlap selection
argmax IoU(mask, box)
mask set
Repairs (optional)
growth; relation operator